# Learning image transformations via convolutional neural networks: a review

Nicholas J. Tustison[1], Brian B. Avants[1], James C. Gee[2],

[1]Department of Radiology and Medical Imaging, University of Virginia, Charlottesville, VA

[2]Department of Radiology, University of Pennsylvania, Philadelphia, PA

Corresponding author:

Nicholas J. Tustison

ntustison@virginia.edu

# Abstract

Recent methodological innovations in deep learning and associated advancements in computational hardware have significantly impacted the various core subfields of quantitative medical image analysis. The generalizability, computational efficiency, and open-source availability of deep learning algorithms, particularly those utlizing convolutional neural networks, have produced paradigm shifts within the field. This impact is evident from topical prevalance in the literature, conference and workshop themes, and winning methodologies in relevant competitions. In this work, we review the various state-of-the-art, fully convolutional network approaches to learning and predicting image transformations. Although of primary importance within the quantitative imaging domain, image registration algorithmic development, in the context of these deep learning strategies, has received comparatively less attention than its counterparts (e.g., image segmentation). Nevertheless, significant inroads have been made and presented in various research venues. We contextualize these contributions within the broader scope of deep learning advancements and, in so doing, attempt to facilitate the leveraging and further development of such techniques within the medical imaging research community.



**Key words:** deep learning, diffeomorphisms, image registration, spatial normalization

# Introduction

Determining the spatial correspondence between imaging domains is frequently a critical component in quantitative image analysis workflows. The evolution of image registration theoretical and technological development has led to increasingly high quality transformational mappings that have significantly improved performance in related processing tasks (e.g., image segmentation via joint label fusion [1]) and imaging-based statistical analyses (e.g., sparse canonical correlation analysis [2]). Several reviews [3–8] have charted this chronology and provided insight into related issues such as algorithmic classification, available implementations, evaluation strategies, and speculation concerning future directions of the field. While prescient in many respects, speculation vis-à-vis deep learning was somewhat limited due to its sudden explosion in popularity and research focus.

The foundational concepts that form the basis for contemporary deep learning studies date back decades (e.g., [9]). Since this early seminal work, major developmental milestones include the *Neocognitron*, an early neural network for character recognition [10], and convolutional neural networks (CNNs or ConvNets) utilized in speech [11] and visual signal processing [12], largely inspired by the visual cell types of the feline visual cortex [13]. The major elements of CNNs are localized connectivity, convolutions, and subsampling (or "pooling") [14]. Furthermore, it is the deep, or hidden, layering that characterizes modern CNNs and is the reason for the extreme performance gains seen with modern architectures. Such architectures are made computationally tractable with gradient-based optimization using backpropagation (first performed in [12]) and the advent of GPU-based hardware [14]. An illustration of a bare-bones CNN configuration is provided in Figure 1 which illustrates the core components of convolution and max pooling. Structural innovations are built upon novel arrangements of these core (and other) network components and the connectivities between them.

A key event in the widespread adoption of CNNs was the 2012 ImageNet Large Scale Visual Recognition Challenge for object classification [15]. The winning entry, a CNN-based architecture colloquially known as *AlexNet* [16], reduced the error rate by almost half over other entries. The following years' competitions were dominated by CNN variants such as VGG [17], GoogLeNet [18], and ResNet [19]. Additional competition outlets including conference-based venues (e.g., NeurIPS) and

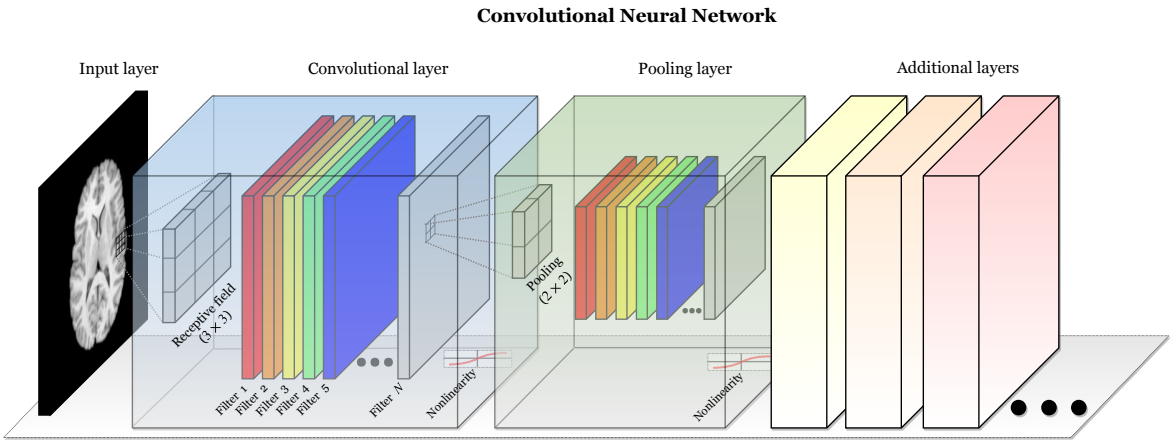


**Figure 1:** The basic elements of the prop convolutional neural network. The convolutional layer comprises several filters which are optimized in terms of their responses to various features found in the input layer. Pooling is used to extract salient features and reduce computational complexity and passed on to subsequent layers.

community-based platforms, such as Kaggle[1], continue to highlight the salience of CNNs as paradigmatic solutions to computational problems. This is in addition to the vast number of formal research reports discussed in the same conferences and published in dedicated journals.

- We need to somehow tie in the reviews [**???**;Schmidhuber:2015aa]
- Uptake in the medical imaging community. Used for such things as image segmentation (U-net). Reviews specific to medical imaging [21–26]
- Early work in medical image registration with the GPU focused on interfacing wiht the hardware directly [27]. One of the review papers listed this as well.

## Spatial transformer networks

In 2015 Jaderberg and his fellow co-authors described a powerful new module, known as the spatial transformer network (STN) [28], which figures prominently in many of the image registration approaches that we review below. Although concepotually relatively straightforward, the obvious influence reflected in recent work will most likely continue in future research which is why we review

[1]Following the 2017 ImageNet challenge, in which the vast majority of teams surpassed the 5% classification error rate threshold, the ImageNet organizers ceded management to the Kaggle community which maintains a running performance assessment in ostensible perpetuity [20].

this important network component.

Generally, STNs enhance CNNs by permiting a flexibility which allows for an explicit spatial invariance that goes beyond the implicitly limited translational invariance associated with the architecture's pooling layers. In many image-based tasks (e.g., localization or segmentation), designing an algorithm that can account for possible pose or geometric variation of the object(s) of interest within the image is crucial for maximizing performance. The STN is a fully differentiable layer which can be inserted anywhere in the CNN to learn the parameters of the transformation of the input feature map (not necessarily an image) which renders the output in such a way to optimize the network based on the specified loss function. The added flexibility and the fact that there is no manual supervision or special handling required makes this module an essential addition for any CNN-based toolkit.

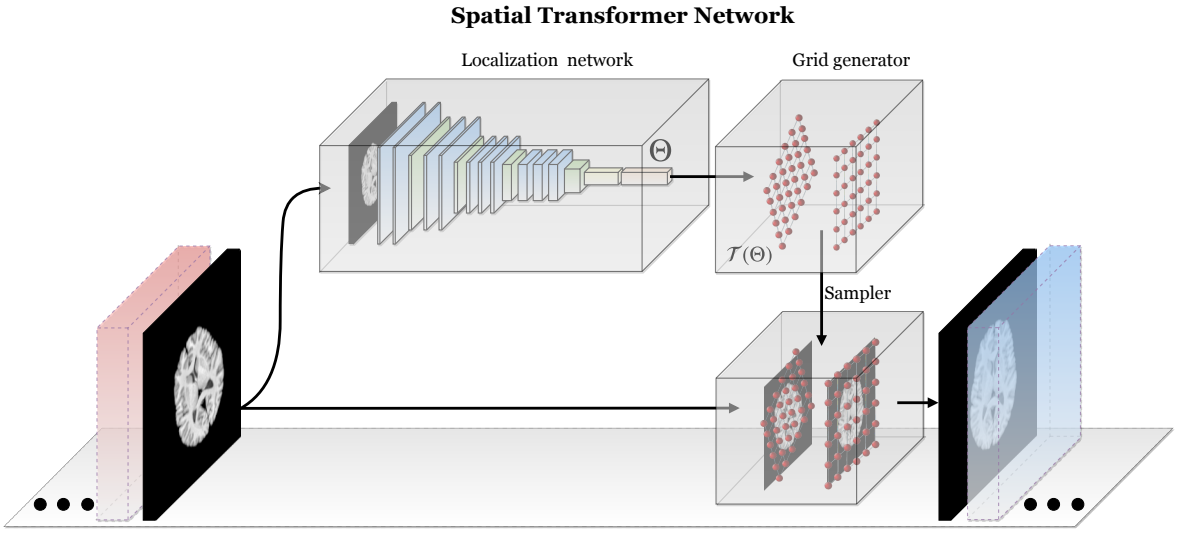


**Figure 2:** Diagrammatic illustration of the spatial transformer network. The STN can be placed anywhere within a CNN to provide spatial invariance for the input feature map. Core components include the localization network used to learn/predict the parameters which transform the input feature map. The transformed output feature map is generated with the grid generator and sampler.

An STN comprises three principal components: 1) a localization network, 2) a grid generator, and 3) a sampler (see Figure 3). The localization network uses the input feature map to learn/regress the transformation parameters which optimize a specified loss function. In many examples provided, this amounts to transforming the input feature map to a quasi-canonical configuration to facilitate, for example, classification. The actual architecture of the localization network is fairly flexible and any conventional architecture, such as a fully connected network (FCN), is suitable as long as the

output maps to the continuous estimate of the transformation parameters. These transformation parameters are then applied to the output of the grid generator which are simply the regular coordinates of the input image (or some normalized version thereof). The sampler, or interpolator, is used to map the transformed input feature map to the coordinates of the output feature map.

## Inverse compositional transformer networks

[29] Inspired by the IC-LK algorithm, we advocate an improved extension to the STN framework that (a) propagates warp parameters, rather than image intensities

## Diffeomorphic transformer networks

Although discussion of transform generalizability was included in the original STN paper [28], discussion was limited to affine, attention (scaling + translation), and thin-plate spline transforms which all fill the requirements of differentiability. This was later extended to encompass a diffeomorphic transformer network (DTN) [30] based on continuous piecewise affine-based (CPAB) transformations [31]. **This section needs to be expanded**.

## Quicksilver

- Quicksilver uses a patch-based approach to predict the "momentum" of the LDDMM framework in determinining the correspondence relationship between fixed and moving image patch pairs.
  - The large deformation diffeomorphic metric mappings (LDDMM) framework for image matching stems from the theoretical foundations underlying diffeomorphic *flows* ([32–34]), or transformations which are differentiable with differentiable inverses. The collection of *paths* between two images which describe the possible mappings between images is described by the
    functional (equation here) XX.

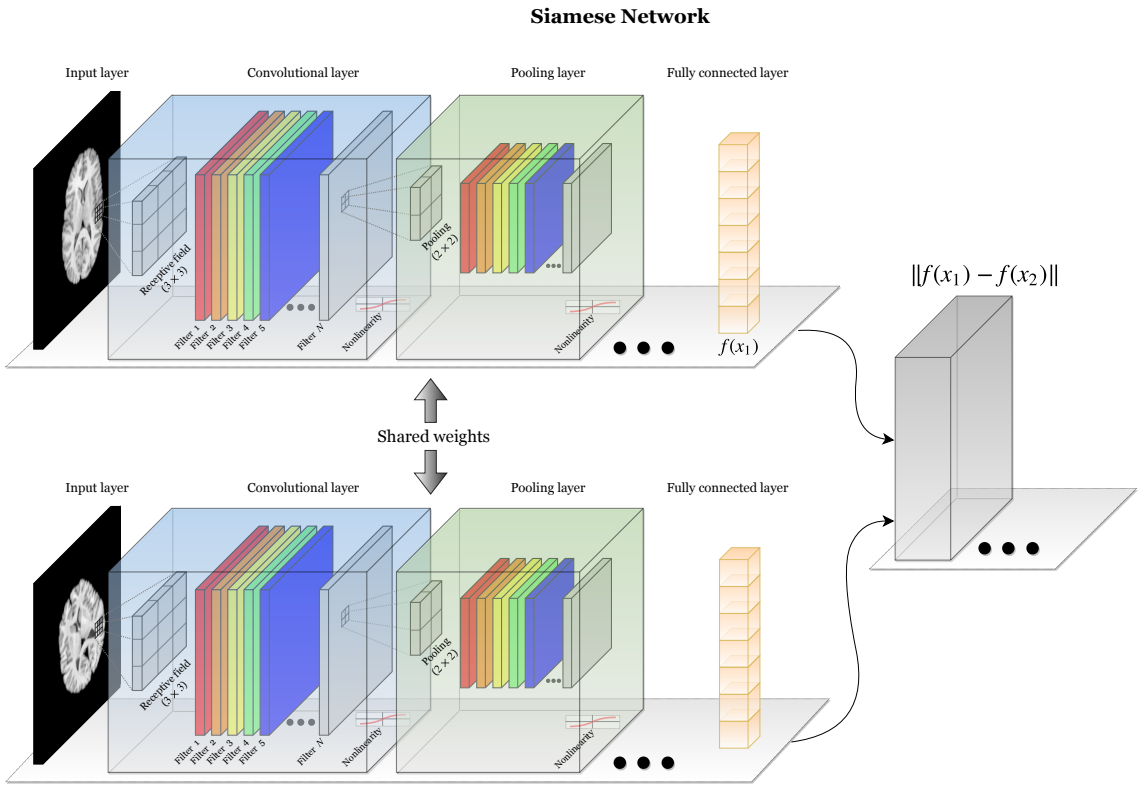


**Figure 3:** Diagrammatic illustration of the spatial transformer network.

- Computationally efficient algorithms for solving such diffeomorphic flows have been subsequently proposed (e.g., [35–40]).
- The Euler-Lagrange equation for the functional XX can be written as a system of equations which incorporates a "momentum" term, $m_t = Jac\left(\phi_{t,1}^v\right)(I \circ \phi_{t,0} - J \circ \phi_{t,1})$. Several algorithms (e.g., [35, 39]) take advantage of the fact that the initial momentum, $m_0$, completely encodes the optimum path describing the transform between images $I$ and $J$.
- Following [41], the initial momentum is calculated from the equation $m(x,0) = \alpha(x,0)\nabla I_0$ where $\alpha$ is a time-varying scalar field function and $\nabla I_0$ is the spatial derivative of the image $I$.
- In costrast to scalar formulations of the momenta (e.g., [39]), a vector formulation is proposed in [42], to avoid potential confounding effects of noise in the image gradient (i.e., better behaved numerically).
- The prediction network is used to estimate the vector momentum

- Network architecture
  - Eencoder/decoder structure for performing voxelwise image regression.
  - There are two separate encoders—one for the fixed image and one for the moving image.
  -
  - $15 \times 15 \times 15$ patches selected at stride = 14.
- Thus, Quicksilver is a supervised approach wherein it uses multiple image pairs and the corresponding ground-truth estimations of the momentum scalar maps to perform prediction of momentum on unseen image pairs (but patch-based).
- Written in pytorch? (torch (facebook) in python) and available at
- Misc. notes
  - Initial momentum is not generally smooth so smoothing is applied after prediction
  - Training data employed regular LDDMM shooting results using PyCA (page 385)
  - They also use dropout layers to induce a probabilistic network
  - They also train a correction network (Section 2.3). THe prediction —> correction networks is diagrammed in Figure 3.
  - The PyCA LDDMM, prediction, and prediction + correction are included in the evaluation.
  - Preprocessing:
    * Skull-striping (FreeSurfer and AutoSeg)
    * Initial affine registration — NiftyReg

# References


1. Iglesias, J. E. and Sabuncu, M. R. "**Multi-Atlas Segmentation of Biomedical Images: A Survey**" *Med Image Anal* 24, no. 1 (2015): 205–219. doi:10.1016/j.media.2015.06.012

2. Avants, B. B., Cook, P. A., Ungar, L., Gee, J. C., and Grossman, M. "**Dementia Induces Correlated Reductions in White Matter Integrity and Cortical Thickness: A Multivariate Neuroimaging Study with Sparse Canonical Correlation Analysis**" *Neuroimage* 50, no. 3 (2010): 1004–16. doi:10.1016/j.neuroimage.2010.01.041

3. Brown, L. G. "**A Survey of Image Registration Techniques**" *ACM Comput. Surv.* 24, no. 4 (1992): 325–376. doi:10.1145/146370.146374, Available at http://doi.acm.org/10.1145/146370.146374

4. Maintz, J. B. and Viergever, M. A. "**A Survey of Medical Image Registration**" *Med Image Anal* 2, no. 1 (1998): 1–36.

5. Pluim, J. P. W., Maintz, J. B. A., and Viergever, M. A. "**Mutual-Information-Based Registration of Medical Images: A Survey**" *IEEE Trans Med Imaging* 22, no. 8 (2003): 986–1004. doi:10.1109/TMI.2003.815867

6. Gholipour, A., Kehtarnavaz, N., Briggs, R., Devous, M., and Gopinath, K. "**Brain Functional Localization: A Survey of Image Registration Techniques**" *IEEE Trans Med Imaging* 26, no. 4 (2007): 427–51. doi:10.1109/TMI.2007.892508

7. Viergever, M. A., Maintz, J. B. A., Klein, S., Murphy, K., Staring, M., and Pluim, J. P. W. "**A Survey of Medical Image Registration - Under Review**" *Med Image Anal* 33, (2016): 140–144. doi:10.1016/j.media.2016.06.030

8. Keszei, A. P., Berkels, B., and Deserno, T. M. "**Survey of Non-Rigid Registration Tools in Medicine**" *J Digit Imaging* 30, no. 1 (2017): 102–116. doi:10.1007/s10278-016-9915-8

9. Ivakhnenko, A. G. "**Polynomial Theory of Complex Systems**" *IEEE Transactions on Systems, Man, and Cybernetics* SMC-1, no. 4 (1971): 364–378.

10. Fukushima, K. "**Neocognitron: A Self Organizing Neural Network Model for a Mecha-**

**nism of Pattern Recognition Unaffected by Shift in Position**" *Biol Cybern* 36, no. 4 (1980): 193–202.

11. Waibel, A. "**Phoneme Recognition Using Time-Delay Neural Networks**" *Meeting of the institute of electrical, information and communication engineers (ieice).* (1987):

12. LeCun, Y., Boser, B., Denker, J. S., Henderson, D., Howard, R. E., Hubbard, W., and Jacke, L. D. "**Backpropagation Applied to Handwritten Zip Code Recognition**" *Neural Computation* 1, no. 4 (1989): 541–551.

13. Hubel, D. H. and Wiesel, T. N. "**Receptive Fields, Binocular Interaction and Functional Architecture in the Cat's Visual Cortex**" *J Physiol* 160, (1962): 106–54.

14. LeCun, Y., Bengio, Y., and Hinton, G. "**Deep Learning**" *Nature* 521, no. 7553 (2015): 436–44. doi:10.1038/nature14539

15. Russakovsky, O., Deng, J., Su, H., Krause, J., Satheesh, S., Ma, S., Huang, Z., Karpathy, A., Khosla, A., Bernstein, M., Berg, A. C., and Fei-Fei, L. "**ImageNet Large Scale Visual Recognition Challenge**" *International Journal of Computer Vision* 115, no. 3 (2015): 211–252.

16. Krizhevsky, A., Sutskever, I., and Hinton, G. E. "**ImageNet Classification with Deep Convolutional Neural Networks**" *Proceedings of the 25th international conference on neural information processing systems - volume 1* (2012): 1097–1105. Available at http://dl.acm.org/citation.cfm?id=2999134.2999257

17. Simonyan, K. and Zisserman, A. "**Very Deep Convolutional Networks for Large-Scale Image Recognition**" *CoRR* abs/1409.1556, (2014): Available at http://arxiv.org/abs/1409.1556

18. Szegedy, C., Vanhoucke, V., Ioffe, S., Shlens, J., and Wojna, Z. "**Rethinking the Inception Architecture for Computer Vision**" *CoRR* abs/1512.00567, (2015): Available at http://arxiv.org/abs/1512.00567

19. He, K., Zhang, X., Ren, S., and Sun, J. "**Deep Residual Learning for Image Recognition**"

*CoRR* abs/1512.03385, (2015): Available at http://arxiv.org/abs/1512.03385

20. Available at https://www.kaggle.com/c/imagenet-object-localization-challenge

21. Greenspan, H., Ginneken, B. V., and Summers, R. M. "**Deep Learning in Medical Imaging: Overview and Future Promise of an Exciting New Technique**" *IEEE Trans Med Imaging* 35, no. 5 (2016): 1153–1159.

22. Suzuki, K. "**Overview of Deep Learning in Medical Imaging**" *Radiol Phys Technol* 10, no. 3 (2017): 257–273. doi:10.1007/s12194-017-0406-5

23. Shen, D., Wu, G., and Suk, H.-I. "**Deep Learning in Medical Image Analysis**" *Annu Rev Biomed Eng* 19, (2017): 221–248. doi:10.1146/annurev-bioeng-071516-044442

24. Litjens, G., Kooi, T., Bejnordi, B. E., Setio, A. A. A., Ciompi, F., Ghafoorian, M., Laak, J. A. W. M. van der, Ginneken, B. van, and Sánchez, C. I. "**A Survey on Deep Learning in Medical Image Analysis**" *Med Image Anal* 42, (2017): 60–88. doi:10.1016/j.media.2017.07.005

25. Ker, J., Wang, L., Rao, J., and Lim, T. "**Deep Learning Applications in Medical Image Analysis**" *IEEE Access* 6, (2018): 9375–9389.

26. Biswas, M., Kuppili, V., Saba, L., Edla, D. R., Suri, H. S., Cuadrado-Godia, E., Laird, J. R., Marinhoe, R. T., Sanches, J. M., Nicolaides, A., and Suri, J. S. "**State-of-the-Art Review on Deep Learning in Medical Imaging**" *Front Biosci (Landmark Ed)* 24, (2019): 392–426.

27. Shams, R., Sadeghi, P., Kennedy, R. A., and Hartley, R. I. "**A Survey of Medical Image Registration on Multicore and the Gpu**" *IEEE Signal Process Mag* 27, no. 2 (2010): 50–60.

28. Jaderberg, M., Simonyan, K., Zisserman, A., and Kavukcuoglu, K. "**Spatial Transformer Networks**" (2015):

29. Lin, C.-H. and Lucey, S. "**Inverse Compositional Spatial Transformer Networks**" *IEEE conference on computer vision and pattern recognition* (2017):

30. Detlefsen, N. S., Freifeld, O., and Hauberg, S. "**Deep Diffeomorphic Transformer Networks**" *Proceedings of the ieee conference on computer vision and pattern recognition conference*

*on computer vision and pattern recognition* (2018):

31. Freifeld, O., Hauberg, S., Batmanghelich, K., and Fisher, J. W. "**Transformations Based on Continuous Piecewise-Affine Velocity Fields**" *IEEE Trans Pattern Anal Mach Intell* 39, no. 12 (2017): 2496–2509. doi:10.1109/TPAMI.2016.2646685

32. Trouvé, A. "**Diffeomorphic Groups and Pattern Matching in Image Analysis**" *Int. J. Computer Vision* 28, (1995): 213–221.

33. Christensen, G. E., Rabbitt, R. D., and Miller, M. I. "**Deformable Templates Using Large Deformation Kinematics**" *IEEE Trans Image Process* 5, no. 10 (1996): 1435–47. doi:10.1109/83.536892

34. Dupuis, P., Grenander, U., and Miller, M. I. "**Variational Problems on Flows of Diffeomorphisms for Image Matching**" *Quarterly of Applied Mathematics* LVI, (1998): 587–600.

35. Beg, M. F., Miller, M. I., Trouvé, A., and Younes, L. "**Computing Large Deformation Metric Mappings via Geodesic Flows of Diffeomorphisms**" *International Journal of Computer Vision* 61, no. 2 (2004): 139–157.

36. Ashburner, J. "**A Fast Diffeomorphic Image Registration Algorithm**" *Neuroimage* 38, no. 1 (2007): 95–113. doi:10.1016/j.neuroimage.2007.07.007

37. Avants, B. B., Epstein, C. L., Grossman, M., and Gee, J. C. "**Symmetric Diffeomorphic Image Registration with Cross-Correlation: Evaluating Automated Labeling of Elderly and Neurodegenerative Brain**" *Med Image Anal* 12, no. 1 (2008): 26–41. doi:10.1016/j.media.2007.06.004

38. Vercauteren, T., Pennec, X., Perchant, A., and Ayache, N. "**Diffeomorphic Demons: Efficient Non-Parametric Image Registration**" *Neuroimage* 45, no. 1 Suppl (2009): S61–72. doi:10.1016/j.neuroimage.2008.10.040

39. Vialard, F.-X., Risser, L., Rueckert, D., and Cotter, C. J. "**Diffeomorphic 3d Image Registration via Geodesic Shooting Using an Efficient Adjoint Calculation**" *Int J Comput Vis* 97, (2012): 229–241.

40. Zhang, M., Liao, R., Dalca, A. V., Turk, E. A., Luo, J., Grant, P. E., and Golland, P. "**Frequency**

**Diffeomorphisms for Efficient Image Registration**" *Inf Process Med Imaging* 10265, (2017): 559–570. doi:10.1007/978-3-319-59050-9_44

41. Miller, M. I., Trouvé, A., and Younes, L. "**Geodesic Shooting for Computational Anatomy**" *J Math Imaging Vis* 24, no. 2 (2006): 209–228. doi:10.1007/s10851-005-3624-0

42. Singh, N., Hinkle, J., Joshi, S., and Fletcher, P. T. "**A Vector Momenta Formulation of Diffeomorphisms for Improved Geodesic Regression and Atlas Construction**" *Proc IEEE Int Symp Biomed Imaging* 2013, (2013): 1219–1222. doi:10.1109/ISBI.2013.6556700